

```
In [1]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import statsmodels.api as sm
from statsmodels.api import OLS, add_constant

from statsmodels.tsa.statespace.sarimax import SARIMAX

plt.rcParams["figure.figsize"] = (16, 9) # Figure size and width
```

```
In [2]: #Omitted Variable Bias Simulation
#Data-generating process

alpha = 0.0          # intercept
beta  = 0.8          # true coefficient on X
delta = 1.2          # effect of omitted Z
rho   = 0.7          # corr(X,Z)
sigma_eps = 0.8      # SD of epsilon
```

```
In [3]: # Monte-Carlo settings

sample_sizes = [50, 300, 2000] # small, medium, large n
reps = 800                     # repetitions
selected_n = 300               # size for histogram plot
np.random.seed(4242)          # reproducibility
```

```
In [4]: #Simulation Loop

summary_rows = []
beta_full_sel = None
beta_omit_sel = None

for n in sample_sizes:
    beta_full = np.empty(reps)
    beta_omit = np.empty(reps)
    for r in range(reps):
        # Generate correlated X and Z
        X = np.random.normal(0, 1, n)
        u = np.random.normal(0, 1, n)
        Z = rho * X + np.sqrt(1 - rho**2) * u

        # Outcome
        eps = np.random.normal(0, sigma_eps, n)
        Y = alpha + beta * X + delta * Z + eps

        # Full model:  $Y \sim X + Z$ 
        X_full = add_constant(np.column_stack((X, Z)))
        beta_full[r] = OLS(Y, X_full).fit().params[1]

        # Reduced model:  $Y \sim X$  (omit Z)
        X_red = add_constant(X)
        beta_omit[r] = OLS(Y, X_red).fit().params[1]

    if n == selected_n:
        beta_full_sel = beta_full
        beta_omit_sel = beta_omit
```

```

summary_rows.append({
    "n": n,
    "mean_beta_full": beta_full.mean(),
    "sd_beta_full": beta_full.std(ddof=1),
    "bias_beta_full": beta_full.mean() - beta,
    "mean_beta_omit": beta_omit.mean(),
    "sd_beta_omit": beta_omit.std(ddof=1),
    "bias_beta_omit": beta_omit.mean() - beta
})

summary = pd.DataFrame(summary_rows)
print("\nMonte Carlo Summary\n", summary)

```

Monte Carlo Summary

	n	mean_beta_full	sd_beta_full	bias_beta_full	mean_beta_omit \
0	50	0.798122	0.166567	-0.001878	1.637012
1	300	0.801978	0.063730	0.001978	1.642243
2	2000	0.800746	0.024714	0.000746	1.641987

	sd_beta_omit	bias_beta_omit
0	0.172766	0.837012
1	0.065205	0.842243
2	0.025419	0.841987

In [5]: *#Theoretical probability Limit*

```

plim_omit = beta + delta * rho
print(f"\nTheoretical plim of  $\beta_{\text{hat}}$  (omit Z) = {plim_omit:.3f}")

```

Theoretical plim of β_{hat} (omit Z) = 1.640

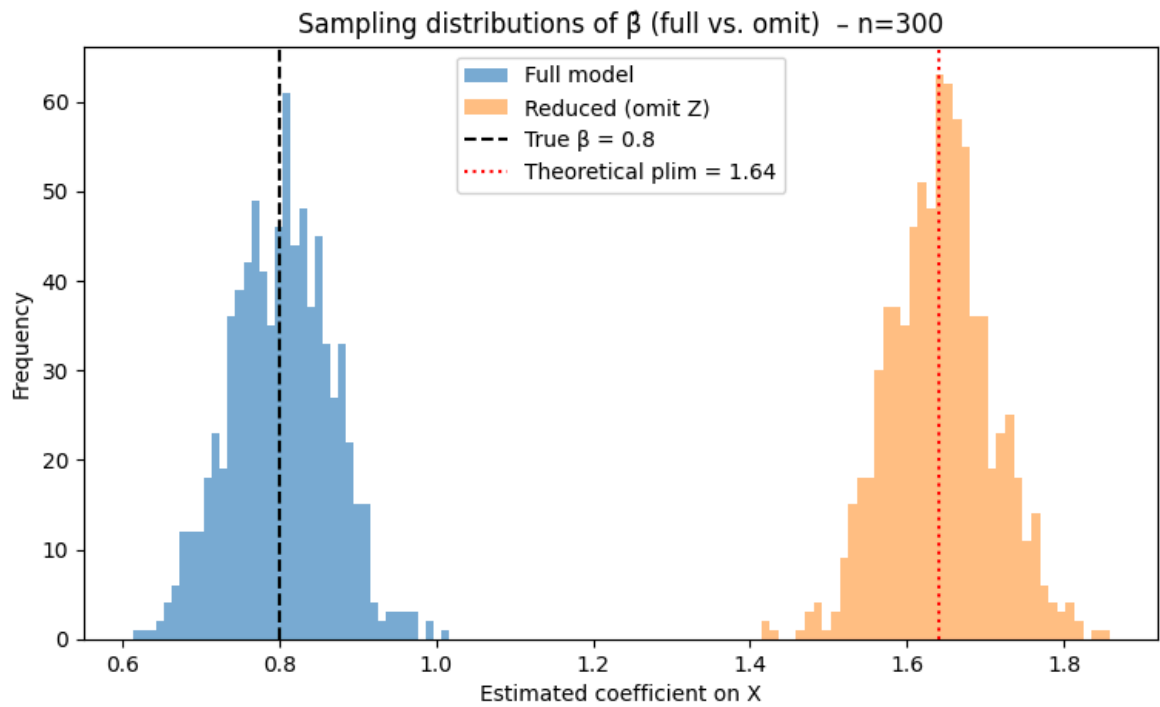
In [6]: *#Plots*

Histogram of sampling distributions at n=300

```

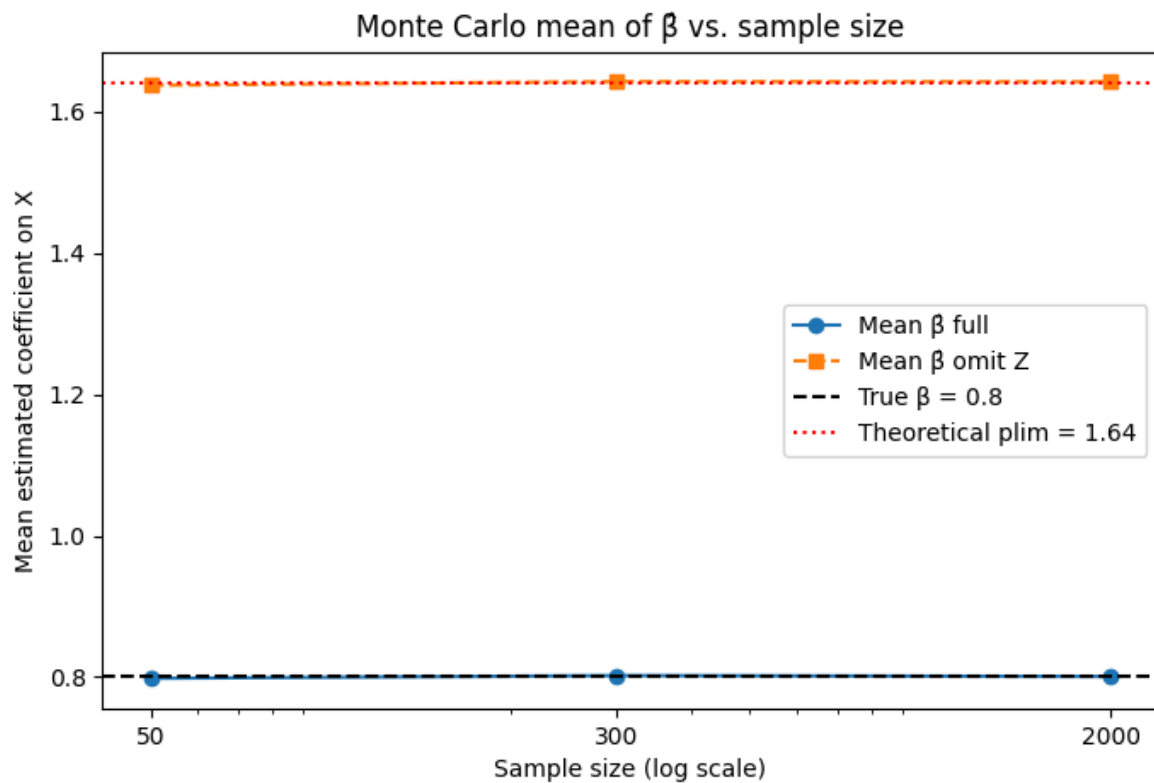
plt.figure(figsize=(9,5))
plt.hist(beta_full_sel, bins=40, alpha=0.6, label="Full model")
plt.hist(beta_omit_sel, bins=40, alpha=0.5, label="Reduced (omit Z)")
plt.axvline(beta, color="black", ls="--", label=f"True  $\beta$  = {beta}")
plt.axvline(plim_omit, color="red", ls=":", label=f"Theoretical plim = {plim_omit}")
plt.title("Sampling distributions of  $\hat{\beta}$  (full vs. omit) - n=300")
plt.xlabel("Estimated coefficient on X")
plt.ylabel("Frequency")
plt.legend()
plt.show()

```



In [7]: *# Mean estimate vs. sample size*

```
plt.figure(figsize=(8,5))
plt.plot(summary["n"], summary["mean_beta_full"], "o-", label="Mean  $\hat{\beta}^{\text{full}}$ ")
plt.plot(summary["n"], summary["mean_beta_omit"], "s--", label="Mean  $\hat{\beta}^{\text{omit Z}}$ ")
plt.axhline(beta, color="black", ls="--", label=f"True  $\beta = \{beta\}$ ")
plt.axhline(plim_omit, color="red", ls=":", label=f"Theoretical plim = {plim_omit}")
plt.xscale("log")
plt.xticks(summary["n"], summary["n"])
plt.title("Monte Carlo mean of  $\hat{\beta}$  vs. sample size")
plt.xlabel("Sample size (log scale)")
plt.ylabel("Mean estimated coefficient on X")
plt.legend()
plt.show()
```



#Summary The coefficient on X for the full model is unbiased; around the true value $\beta=0.8$. The coefficient is larger in the reduced model; about 1.64 which is equal to the theoretical probability limit $= \beta + \delta\rho = 0.8 + 1.2 \times 0.7 = 1.64$. The sampling variance shrinks in the full model as n grows. The sampling variance also shrinks as n increases in the reduced model and the centre stays at 1.64. In conclusion, increasing the sample size does not eliminate omitted-variable bias. It only reduces random sampling error.